

1 Definitions

- We represent probabilistic models using factor graphs, denoted (θ, F) of variables $\theta = (\theta_1, \dots, \theta_N)$ and factors $F = (F_1, \dots, F_k)$.

Θ denotes the space of variable values. $\theta \in \Theta$ contains all parameters, latent variables, and observations (which are fixed).

F represents all functions, operations, and constraints that the posterior distribution factorizes into from a Bayesian perspective including priors and likelihoods (in directed graphs). Given data, the unnormalized posterior distribution can then be expressed as

$$\prod_k F_k(\theta)$$

- Symmetry $\sigma : \Theta \rightarrow \Theta$ is a measurable function with a measurable inverse satisfying

$$\prod_k F_k(\theta) \propto \prod_k F_k(\sigma(\theta))$$

where the product is taken over non-prior terms. Symmetries must also fix the components of θ that correspond to observed variables.

- Local symmetry is a symmetry σ that satisfies

$$F_k(\theta) \propto F_k(\sigma(\theta))$$

for all non-prior factors F_k .

1.1 Equivalence class

- A factor graph with equivalence relation Ξ , denoted $\mathcal{F} = (\theta, F, \Xi)$, induces a partition of Θ into equivalence classes $[\cdot]$, with the equivalence relation Ξ that there exists a symmetry between any two members of an equivalence class. Ξ specifies the type of symmetry.

Denote $[\theta]$ to be the equivalence class of θ .

Denote Θ/Ξ to be the set of all equivalence classes in Θ w.r.t Ξ , called the quotient set of Θ by Ξ .

(TODO: local symmetry?)

- (TODO: Specific to breaking symmetry post-inference?)

A symmetry detector $\Delta_{\mathcal{F}} : \Theta \rightarrow \Theta/\Xi$ is a measurable function that maps θ to $[\theta]$ w.r.t. Ξ .

- (TODO: Specific to breaking symmetry post-inference?)

A symmetry breaker $\phi_{\mathcal{F}} : \Theta/\Xi \rightarrow \Theta$ is a measurable function that maps $[\theta]$ to θ_0 w.r.t. Ξ where θ_0 is a representative of $[\theta]$.

- (TODO: Specific to breaking symmetry post-inference?)

An automatic symmetry breaker $\Phi_{\mathcal{F}} : \Theta \rightarrow \Theta$ is the composition $\Phi_{\mathcal{F}} = \phi_{\mathcal{F}} \circ \Delta_{\mathcal{F}}$ of a symmetry breaker with the corresponding symmetry detector.

(TODO)

An automatic symmetry breaker $\Phi_{\mathcal{F}} : \Theta \rightarrow \Theta$ is the composition $\Phi_{\mathcal{F}} = \phi_{\mathcal{F}} \circ \Delta_{\mathcal{F}}$ of a symmetry breaker with the corresponding symmetry detector that satisfies the following:

1. For ASB for two types of symmetries Φ_1, Φ_2 , $\Phi_1 \circ \Phi_2(\theta) = \Phi_2 \circ \Phi_1(\theta)$

Has the following properties:

1. $\Phi \circ \Phi(\theta) = \Phi(\theta)$
2. $\forall \theta_* \in [\theta], \Phi(\theta_*) = \Phi(\theta_0)$ where θ_0 representative of $[\theta]$

- A symmetry broken posterior distribution is a posterior distribution that has the (unnormalized) form

$$\prod_k F_k(\Phi_{\mathcal{F}_M} \circ \dots \circ \Phi_{\mathcal{F}_1}(\theta))$$

where $\mathcal{F}_1, \dots, \mathcal{F}_M$ have equivalence relations Ξ_1, \dots, Ξ_M corresponding to different types of symmetries equipped, respectively.

2 Breaking symmetries post-inference

Context: Have a posterior distribution $\prod_k F_k(\theta)$ where inference has been performed (θ has/can be identified)

Problem: Nonidentifiability. There exists a (local) symmetry σ s.t. $\prod_k F_k(\theta) \propto \prod_k F_k(\sigma(\theta))$ or for local, $F_k(\theta) \propto F_k(\sigma(\theta))$

When: Inference algorithm performance is not an issue in the presence of a symmetry.

Solution: Use the symmetry broken posterior distribution

2.1 Scaling symmetries

Scaling symmetry is a symmetry σ which multiplies θ pointwise by a vector

$$v = (r_1, \dots, r_N) = (e^{d_1}, \dots, e^{d_N})$$

where $r_n \in \mathbb{R}_+$ and $d_n = \log r_n$.

- Let Σ_v be the diagonal matrix with v on the diagonal. Then $\sigma(\theta) = \Sigma_v \theta$.
- Let $d = (d_1, \dots, d_N)$. Let $v_d = (e^{d_1}, \dots, e^{d_N})$. Then the scaling symmetries of the model are $\sigma_d(\theta) = \Sigma_{v_d} \theta$, $d \in \mathcal{N}(C)$.
- Then $\{\Sigma_{v_d} \theta : d \in \mathcal{N}(C)\} \subset \Theta$ is an equivalence class for the model (θ, F) .

3 Notes

- Class of transformations of model parameters $T \subset \{f : \Theta \rightarrow \Theta\}$. Subset of local symmetries contained in T $S_T \subset T$ defined by

$$S_T = \{\sigma \in T | \sigma \text{ is a local symmetry} \} = \bigcap_k \{\sigma \in T | F_k(\theta) \propto F_k(\sigma(\theta))\}$$

- Equivalence class definition

A factor graph (model?) with equivalence relation Ξ , denoted $\mathcal{F} = (\theta, F, \Xi)$, induces a partition of Θ into equivalence classes, with the equivalence relation Ξ that there exists a symmetry between any two members of an equivalence class. Ξ specifies the type of symmetry.

Classes may be singletons if the point is symmetric only to itself.

Denote $[\theta]$ to be the equivalence class of θ .

Denote Θ/Ξ to be the set of all equivalence classes in Θ w.r.t Ξ , called the quotient set of Θ by Ξ .

- A symmetry detector $\Delta_{\mathcal{F}} : \Theta \rightarrow \Theta/\Xi$ is a measurable? function that maps θ to $[\theta]$ w.r.t. Ξ .

Symmetry detection is identifying the equivalence class?

- Definition 1: A symmetry breaker $\phi_{\mathcal{F}} : \Theta/\Xi \rightarrow \Theta$ is a measurable function that maps $[\theta]$ to θ_0 w.r.t. Ξ where θ_0 is a representative of $[\theta]$.

Definition 2: A symmetry breaker $\phi_{\mathcal{F}} : \Theta/\Xi \rightarrow \Theta$ is a measurable function that maps $[\theta]$ to $\theta \in [\theta]$ w.r.t. Ξ .

- An automatic symmetry breaker $\Phi_{\mathcal{F}} : \Theta \rightarrow \Theta$ is the composition $\Phi_{\mathcal{F}} = \phi_{\mathcal{F}} \circ \Delta_{\mathcal{F}}$ of a symmetry breaker with the corresponding symmetry detector.

A symmetry broken posterior distribution is a posterior distribution that has the (unnormalized) form

$$\prod_k F_k(\Phi_{\mathcal{F}_M} \circ \dots \circ \Phi_{\mathcal{F}_1}(\theta))$$

where $\mathcal{F}_1, \dots, \mathcal{F}_M$ have equivalence relations Ξ_1, \dots, Ξ_M corresponding to different types of symmetries equipped, respectively.

Φ is not a symmetry (no measurable inverse)

- Scaling symmetry is a symmetry σ which multiplies θ pointwise by a vector

$$v = (r_1, \dots, r_N) = (e^{d_1}, \dots, e^{d_N})$$

where $r_n \in \mathbb{R}_+$ and $d_n = \log r_n$.

- Case of $r_n \in \mathbb{R}_-$ is covered by a combination of scaling symmetry and sign-flip symmetry.

- Let Σ_v be the diagonal matrix with v on the diagonal. Then $\sigma(\theta) = \Sigma_v \theta$.
- Let matrix C be the matrix of constraints on d_n . The scaling symmetries of the model are the vectors in the null space of C , $\mathcal{N}(C)$.
- Let $d = (d_1, \dots, d_N)$. Let $v_d = (e^{d_1}, \dots, e^{d_N})$. Then the scaling symmetries of the model are $\sigma_d(\theta) = \Sigma_{v_d} \theta$, $d \in \mathcal{N}(C)$.
- Then $\{\Sigma_{v_d} \theta : d \in \mathcal{N}(C)\} \subset \Theta$ is an equivalence class for the model (θ, F) .
- The symmetry detector $\Delta_{\mathcal{F}_S}$ for a scaling symmetry is defined as

$$\Delta_{\mathcal{F}_S}(\theta) = \{\Sigma_{v_d} \theta : d \in \mathcal{N}(C)\}$$

- A symmetry breaker (definition) for a scaling symmetry is

$$\phi_{\mathcal{F}_S}([\theta]) = \begin{cases} \theta & \mathcal{N}(C) = \{\vec{0}\} \\ \Sigma_{v_d} \theta, \|v_d\| = 1 & \mathcal{N}(C) \neq \{\vec{0}\} \end{cases}$$

- A symmetry breaker (definition 1) for a scaling symmetry is

$$\phi_{\mathcal{F}_S}([\theta]) = \Sigma_d \theta$$

Let $\theta_* = \Sigma_{d^*} \theta$ for some $d^* \in \mathcal{N}(C)$.

Then

$$\begin{aligned} \Delta_{\Xi}(\theta_*) &= \{\Sigma_d \theta_* : d \in \mathcal{N}(C)\} \\ &= \{\Sigma_d \Sigma_{d^*} \theta : d \in \mathcal{N}(C)\} \\ &= \{\Sigma_{d+d^*} \theta : d \in \mathcal{N}(C)\} \\ &= \{\Sigma_d \theta : d \in \mathcal{N}(C)\} \end{aligned}$$

where the above follows as $d + d^* \in \mathcal{N}(C)$. Hence Δ_{Ξ} maps members of an equivalence class to the equivalence class.

Symmetry breaker: need to constrain $d \in \mathcal{N}(C)$ to get a unique representative? How to remove dependence on θ ?

- G&H 5.3: latent variable $z_i = X_i \beta + \epsilon_i$, $\epsilon_i \sim N(0, \sigma^2)$, scaling of β doesn't change sign of z_i . Proposed solution is just to fix σ .
- N'98: same problem as above but now covariance matrix. Proposed solution is just to fix $\Sigma_{1,1} = 1$.